

# CS 559: Machine Learning Fundamentals & Applications

## Lecture 4: Linear Regression





## 4.1. Linear Basis Function Models

- 4.1.1. Maximum Likelihood and Least Squares

- 4.1.2. Sequential Learning

- 4.1.3. sklearn Linear Regression Example

## 4.2. Regularization

- 4.2.1. Overfit vs. Underfit

- 4.2.2. Regularized Least Squares

## 4.3. The Bias-Variance Decomposition

## 4.4. Conclusion



## 4.1. Linear Basis Function Models

4.1.1. Maximum Likelihood and Least Squares

4.1.2. Sequential Learning

4.1.3. sklearn Linear Regression Example

## 4.2. Regularization

4.2.1. Overfit vs. Underfit

4.2.2. Regularized Least Squares

## 4.3. The Bias-Variance Decomposition

## 4.4. Conclusion



## 4.1. Linear Basis Function Models

The simplest linear model for regression is called linear regression that is a linear combination of the input variables

$$y(\mathbf{x}, \mathbf{w}) = w_0 + w_1 x_1 + \cdots w_D x_D. \quad (4-1)$$

- Input variable  $\mathbf{x} = (x_1, \dots, x_D)^T$
- Parameters  $\mathbf{w} = (w_0, w_1, \dots, w_D)$
- For a simplicity, Eq. (4-1) above can be expressed as

$$y(\mathbf{x}, \mathbf{w}) = \sum_{i=0}^D w_i x_i = \mathbf{w} \mathbf{x} \quad (4-2)$$

where  $\mathbf{x} = (1, x_1, \dots, x_D)^T$ .



## 4.1. Linear Basis Function Models

If  $y(\mathbf{x}, \mathbf{w})$  is **not linear**, the model will impose significant limitation. Instead, we can make the model be a linear **combination of fixed nonlinear function** of the input variables in the form

$$y(\mathbf{x}, \mathbf{w}) = w_0 + \sum_{j=1}^{M-1} \mathbf{w}_j \phi_j(\mathbf{x}) = \sum_{j=0}^{M-1} w_j \phi_j(\mathbf{x}) = \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x}) \quad (4-3)$$


where  $\phi_j(\mathbf{x})$  are known as **basis function**.

When nonlinear basis function is used,  $y$  is not linear in terms of  $\mathbf{x}$  but is linear in  $\mathbf{w}$ . This linearity in parameters simplifies the model greatly!

## 4.1. Linear Basis Function Models


Some possible choices for the basis functions are

$$\phi_j(x) = \exp \left\{ -\frac{(x - \mu_j)^2}{2\sigma^2} \right\}$$
$$\phi_j(x) = \sigma \left( \frac{x - \mu_j}{s} \right) \text{ where } \sigma(a) = \frac{1}{1 + \exp(-a)}$$




# 4.1. Linear Basis Function Models

$$y = ax^2 + bx + c$$




$$y = w_1 \frac{1}{1 + \exp(-x)} + w_0$$



$$\phi(x) = \frac{1}{1 + \exp(-x)}$$


$$y = w_1 \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right) + w_0$$

$$\phi(x) = \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right)$$



$$y = w_1(x^2 + x) + w_0$$

$$\phi(x) = (x^2 + x)$$





## 4.1.1. Maximum Likelihood and Least Squares

Assume the target variable  $t$  is given by a deterministic function  $y(\mathbf{x}, \mathbf{w})$  with additive Gaussian noise  $\epsilon = \mathcal{N}(0, \beta^{-1})$  where  $\beta$  is the inverse of variance ( $\beta = \sigma^{-2}$ ),

$$t = y(\mathbf{x}, \mathbf{w}) + \epsilon. \quad (4-4)$$

Recall the conditional probability in Gaussian, the likelihood of Eq. (4-4) can be expressed as

$$p(t|x, w, \beta) = \mathcal{N}(t|y(x, w), \beta^{-1}). \quad (4-5)$$

Consider a data set of inputs  $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$  with corresponding target values  $\mathbf{t} = (t_1, \dots, t_N)$ . Eq. (4-5) becomes

$$\begin{aligned} p(\mathbf{t}|\mathbf{X}, \mathbf{w}, \beta) &= \prod_{n=1}^N \mathcal{N}(t_n | \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x}_n), \beta^{-1}) \\ &= \prod_{n=1}^N \frac{\beta^{N/2}}{\sqrt{(2\pi)^N}} \exp\left(-\frac{\beta}{2} (t_n - \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x}_n))^2\right) \end{aligned} \quad (4-6)$$





## 4.1.1. Maximum Likelihood and Least Squares

Assume that data points are i.i.d., then we can write Eq. 4-6 as  $p(\mathbf{t}|\mathbf{X}, \mathbf{w}, \beta) = p(\mathbf{t}|\mathbf{w}, \beta)$ . Then log-likelihood then becomes

$$\begin{aligned}\ln p(\mathbf{t}|\mathbf{w}, \beta) &= \sum_{n=1}^N \ln \mathcal{N}(t_n | \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x}_n), \beta^{-1}) \\ &\quad \uparrow \mathcal{N} = \frac{\beta^{N/2}}{\sqrt{(2\pi)^N}} \exp\left(-\frac{\beta}{2} (t_n - \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x}_n))^2\right) \\ &= \frac{N}{2} \ln \beta - \frac{N}{2} \ln(2\pi) - \frac{\beta}{2} \sum_{n=1}^N \{t_n - \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x}_n)\}^2 \\ &= \frac{N}{2} \ln \beta - \frac{N}{2} \ln(2\pi) - \beta E_D(\mathbf{w})\end{aligned}\tag{4-7}$$

where the sum-square-error function is  $E_D = \frac{1}{2} \sum_{n=1}^N \{t_n - \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x}_n)\}^2$ .



## 4.1.1. Maximum Likelihood and Least Squares

To find  $\mathbf{w}$ , we need to maximize the log-likelihood function that is equivalent to *minimizing*  $E_D(\mathbf{w})$ :

$$\nabla \ln p(\mathbf{t}|\mathbf{w}, \beta) = 0 \rightarrow \nabla E_D(\mathbf{w}) = \frac{\partial}{\partial \mathbf{w}} E_D(\mathbf{w}) = \sum_{n=1}^N \{t_n - \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x}_n)\} \boldsymbol{\phi}(\mathbf{x}_n)^T = 0 \quad (4-8)$$

where  $\beta$  is the constant and therefore, it does not impact the result of  $\mathbf{w}$ .

Solving for  $\mathbf{w}$ ,

$$\mathbf{w}_{ML} = (\boldsymbol{\Phi}^T \boldsymbol{\Phi})^{-1} \boldsymbol{\Phi}^T \mathbf{t} \quad (4-9)$$

which is known as the *normal equations* for the least squares problem. The *design matrix*  $\boldsymbol{\Phi}$  is an  $N \times M$  matrix whose elements are  $\Phi_{nj} = \phi_j(\mathbf{x}_n)$ :

$$\boldsymbol{\Phi} = \begin{pmatrix} \phi_0(\mathbf{x}_1) & \cdots & \phi_{M-1}(\mathbf{x}_1) \\ \vdots & \ddots & \vdots \\ \phi_0(\mathbf{x}_N) & \cdots & \phi_{M-1}(\mathbf{x}_N) \end{pmatrix}.$$

## 4.1.1. Maximum Likelihood and Least Squares

If  $w_0$  is explicit, then

$$E_D = \frac{1}{2} \sum_{n=1}^N \left\{ t_n - w_0 - \sum_{j=1}^{M-1} w_j \phi_j(x_n) \right\}^2 \quad (4-10)$$

and

$$w_0 = \bar{t} - \sum_{j=1}^{M-1} w_j \bar{\phi}_j \quad (4-11)$$

where

$$\bar{t} = \frac{1}{N} \sum_{n=1}^N t_n, \quad \bar{\phi}_j = \frac{1}{N} \sum_{n=1}^N \phi_j(x_n). \quad (4-12)$$



## 4.1.1. Maximum Likelihood and Least Squares

We can also estimate the noise precision parameter  $\beta$ ,

$$\frac{1}{\beta_{ML}} = \frac{1}{N} \sum_{n=1}^N \{t_n - \mathbf{w}_{ML}^T \boldsymbol{\phi}(\mathbf{x}_n)\}^2 \quad (4-12)$$

where  $\sigma_{ML}^2 = \beta_{ML}^{-1}$ .



## 4.1.2. Sequential Learning

If the data set is sufficiently large, it is worth using the *sequential* algorithm, also known as the *online* algorithm. The solutions obtained via MLE (Eq. 4-8) will be underestimated.

The technique of *stochastic gradient descent* (aka *sequential gradient descent*) updates the parameter  $\mathbf{w}$  through the iteration  $\tau$  with an initial set of  $\mathbf{w}^{(0)}$ :

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} - \eta \nabla E_n$$

$$\nabla E_n = \partial E_n / \partial \mathbf{w}$$

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} + \eta (t_n - \mathbf{w}^{(\tau)T} \boldsymbol{\phi}_n) \boldsymbol{\phi}_n \quad (4-13)$$

where  $\eta$  is a learning rate parameter and  $\boldsymbol{\phi}_n = \phi(\mathbf{x}_n)$  and it is  $\eta \ll 1$ .



## 4.1.2. Sequential Learning

### 1. Gradient Descent (GD, bath gradient descent):

- It involves using the entire dataset or training set to compute the gradient to find the optimal solution.
- Our movement towards the optimal solution, which could be the local or global optimal solution, is always direct.
- However, this can become a major challenge when millions of samples are running.

### 2. Stochastic Gradient Descent (SGD):

- The dataset is properly shuffled to avoid pre-existing orders then partitioned into  $m$  examples or an example at a time after reordering.
- This random approximation of the data set removes the computational burden associated with gradient descent while achieving iteration faster and at a lower convergence rate.
- It tends to result in more noise than GD.

### 3. Mini-Batch Gradient Descent: Bridge between GD and SGD.

## 4.1.2. Sequential Learning

Suppose there are  $m$  examples,  $\mathbf{x}$  has  $n$  many features.

- GD: Repeat until the error function converges

---

```
Repeat {  
  for  $j = 1, \dots, n$  {  
     $w_j = w_j - \frac{\eta}{m} \sum_{i=1}^m (y_w(\mathbf{x}^{(i)}) - \mathbf{t}^{(i)}) \mathbf{x}_j^{(i)}$   
  }  
}
```

---

- SGD: Randomly shuffle and repeat until the error function converges

---

```
Repeat {  
  for  $i = 1, \dots, m$  {  
     $w_j = w_j - \eta (y_w(\mathbf{x}^{(i)}) - \mathbf{t}^{(i)}) \mathbf{x}_j^{(i)} \quad \forall j = [0, \dots, n]$   
  }  
}
```

---



## 4.1.2. Sequential Learning

- Mini-Batch GD: Randomly shuffle and repeat the error function until converges
  - Suppose we make 10 batches with 1000 examples.

---

Repeat {

for  $i = 1, 11, 21, \dots, 991$  {

$$w_j = w_j - \frac{\eta}{10} \sum_{k=i}^{i+9} (y_w(x^{(k)}) - t^{(k)}) x_j^{(k)} \quad \forall j = [0, \dots, n]$$

}

}

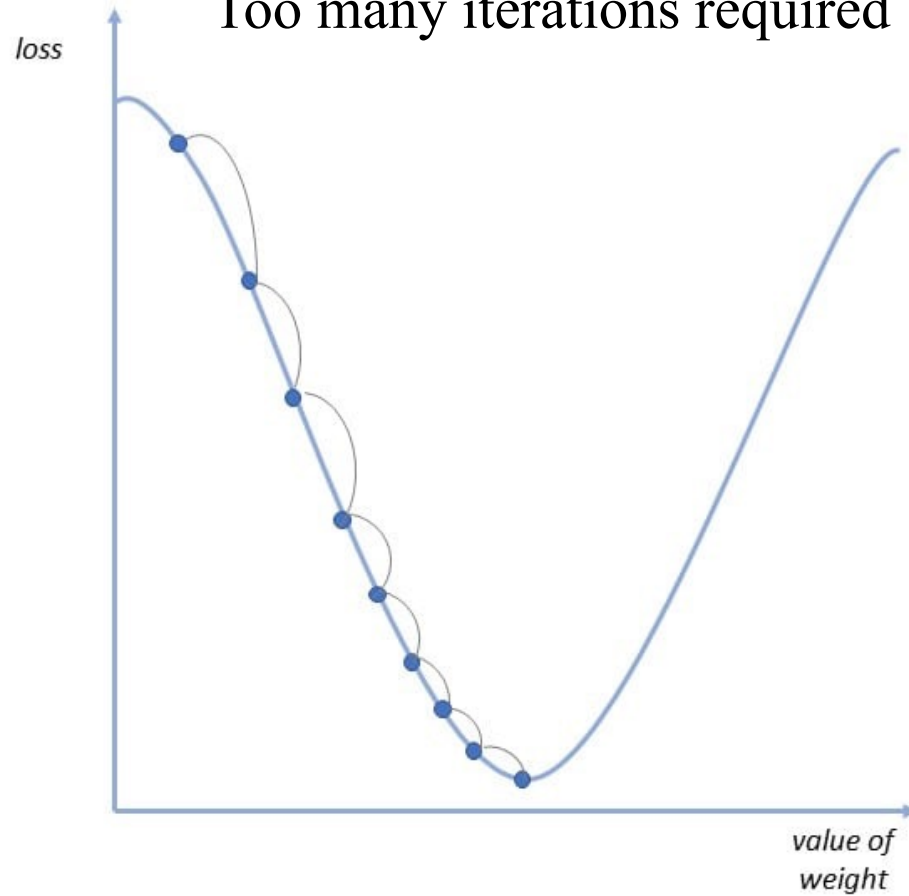
---



## 4.1.2. Sequential Learning


Small Learning Rate

Too many iterations required




Large Learning Rate

May never reach the global/local minimum  
We overshoot the target!



## 4.1.2. Sequential Learning






## 4.1.2. Sequential Learning


```
1 np.random.seed(123)
2 X = 10*np.random.sample(100)-5
3 T = 3*X
4 plt.scatter(X,T,s=1.2)
5 plt.ylabel('T')
6 plt.xlabel('X')
7 plt.show()
```

✓ 0.4s




```
1 beta = 2
2 learning_rate = [1E-3, 1E-2, 0.1, 5E-1]
3 beta_list = []
4 for eta in learning_rate:
5     beta_temp = []
6     for i in range(4):
7         Y = beta*X
8         grad_loss = np.mean(2*(T-beta*X)*-X)
9         beta = beta - eta*grad_loss
10        beta_temp.append(beta)
11    beta_list.append(beta_temp)
```


```
eta=0.001,beta=[2.0120427634766513, 2.023940498801148, 2.035694952513224, 2.0473078501194486]
eta=0.01,beta=[2.162038312390189, 2.2629520564539956, 2.3517129970047645, 2.429784667425353]
eta=0.1,beta=[3.116481505521015, 2.9762055834816294, 3.004860636501196, 2.999007086928208]
eta=0.5,beta=[3.0049857955664416, 2.9749644171916443, 3.1257132183222187, 2.3687459412506553]
```



## 4.1.3. sklearn Linear Regression



## 4.1.3. sklearn Linear Regression



```
1 ols.fit(x_m, y_m)
2 print("beta_1, beta_2: " + str(np.round(ols.coef_, 3)))
3 print("beta_0: " + str(np.round(ols.intercept_, 3)))
4 print("RSS: %.2f" % np.sum((ols.predict(x_m) - y_m) ** 2))
5 print("R^2: %.5f" % ols.score(x_m, y_m))
```

```
beta_1, beta_2: [-6.619  4.436]
beta_0: 2.523
RSS: 356.34
R^2: 0.83938
```

Ordinary Least Square (OLS) is the result of MLE:

If  $y_i = \beta_0 + \beta x_i + \epsilon_i$ ,

$$\beta = \frac{n \sum x_i y_i - \sum x_i \sum y_i}{n \sum x_i^2 - (\sum x_i)^2}, \beta_0 = \bar{y} - \beta \bar{x}.$$

If  $y_i = \sum_{j=1}^n \beta_j X_{ij}$ ,  $(i = 1, \dots, n)$ ,

$$\beta = \frac{X^T y}{(X^T X)^{-1}}$$



## 4.1.3. sklearn Linear Regression

```
1 from sklearn.linear_model import LinearRegression
2 from sklearn.model_selection import train_test_split
3 from sklearn.metrics import mean_squared_error, mean_absolute_error
✓ 0.3s
```

Parameters:

- **fit\_intercept** (*bool, default=True*): If set to False, no intercept will be used in calculations (i.e. data is expected to be centered).
- **Normalize** (*bool, default=False*): This parameter is ignored when **fit\_intercept** is set to False. If True, the regressors X will be normalized before regression. Need to standardize the data with false when the standardized data is used.
- **copy\_X** (*bool, default=True*): If True, X will be copied; else, it may be overwritten.
- **n\_jobs** (*int, default=None*): The number of jobs to use for the computation. This will only provide speedup in case of sufficiently large problems, that is if firstly `n_targets > 1` and secondly X is sparse or if positive is set to True. None means 1 unless in a [joblib.parallel\\_backend](#) context. -1 means using all processors. See [Glossary](#) for more details.
- **positive**: *bool, default=False*. When set to True, forces the coefficients to be positive. This option is only supported for dense arrays.



## 4.1.3. sklearn Linear Regression

Recall the data from lecture 2,

longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value	rooms_per_household	bedrooms_per_room	population_per_household	<1H OCEAN	INLAND	ISLAND	NEAR BAY	NEAR OCEAN
-122.23	37.88	32.672294	8.208606	5.645033	7.675271	5.705034	2.728723	61.003713	3.113149	-2.957260	0.876915	0	0	0	1	0
-122.22	37.86	17.070476	11.413717	8.717972	11.487953	8.974699	2.724071	58.036799	2.848110	-2.822167	0.707407	0	0	0	1	0
-122.26	37.84	33.434943	9.799169	7.954136	10.099133	7.961036	0.797670	52.585885	2.062241	-1.811675	0.671329	0	0	0	1	0
-122.26	37.85	39.491212	8.560402	6.722807	9.055987	6.750997	0.822822	47.360433	2.038908	-1.864239	0.905250	0	0	0	1	0
-122.26	37.84	39.491212	9.597088	7.399371	9.717414	7.430798	0.746034	48.672270	2.506253	-2.271723	0.808116	0	0	0	1	0

```
Y = df['median_house_value']
features = list(set(df.columns.tolist()[3:])-set(['median_house_value']))
print(features)
X = df[features]
X_train, X_test, y_train, y_test = train_test_split(X,Y,test_size=0.3,random_state=42)
```





## 4.1.3. sklearn Linear Regression

```
def lr_model(xT,xt,yT,yt,features):  
    """  
    xT: Original Train X, yT: Original Train Y  
    xt: Original Test X, yt: Original Test Y  
    features: features in train dataset  
    """  
    XT,Xt = xT[features], xt[features]  
    lr = LinearRegression() #Make a model  
    lr.fit(XT,yT) #Fit a train data  
    y_train_pred = lr.predict(XT) #Predict the train data  
    rmse = mean_squared_error(yT,y_train_pred) #Calculate root-mean-squared-error  
    mae = mean_absolute_error(yT,y_train_pred) #Calculate mean-absolute-error  
    rsq = r2_score(yT,y_train_pred) #Calculate R^2  
    print(f'train: rmse={rmse}, mae={mae}, R^2={rsq}')  
    y_test_pred = lr.predict(Xt) #Predict the test data  
    rmse = mean_squared_error(yt,y_test_pred)  
    mae = mean_absolute_error(yt,y_test_pred)  
    rsq = r2_score(yt,y_test_pred) #Calculate R^2  
    print(f'train: rmse={rmse}, mae={mae}, R^2={rsq}')  
    train_res = yT - y_train_pred #Calculate residuals  
    test_res = yt - y_test_pred  
    plt.scatter(yT,train_res,alpha=0.4,label='train') #Plot residuals  
    plt.scatter(yt,test_res,alpha=0.4, label='test')  
    plt.hlines(y=0,xmin=yT.min()-2,xmax=yT.max()+2,linestyles='--',color='r')  
    plt.legend()  
    plt.show()  
    return lr
```



## 4.1.3. sklearn Linear Regression



- Total Sum of Squares (TSS):

$$\sum_i^N (y_i - \bar{y})^2$$

- Residual Sum of Squares (RSS):

$$\sum_i^N (t_i - y_i)^2$$

- R-square:

$$R^2 = 1 - \frac{RSS}{TSS}$$

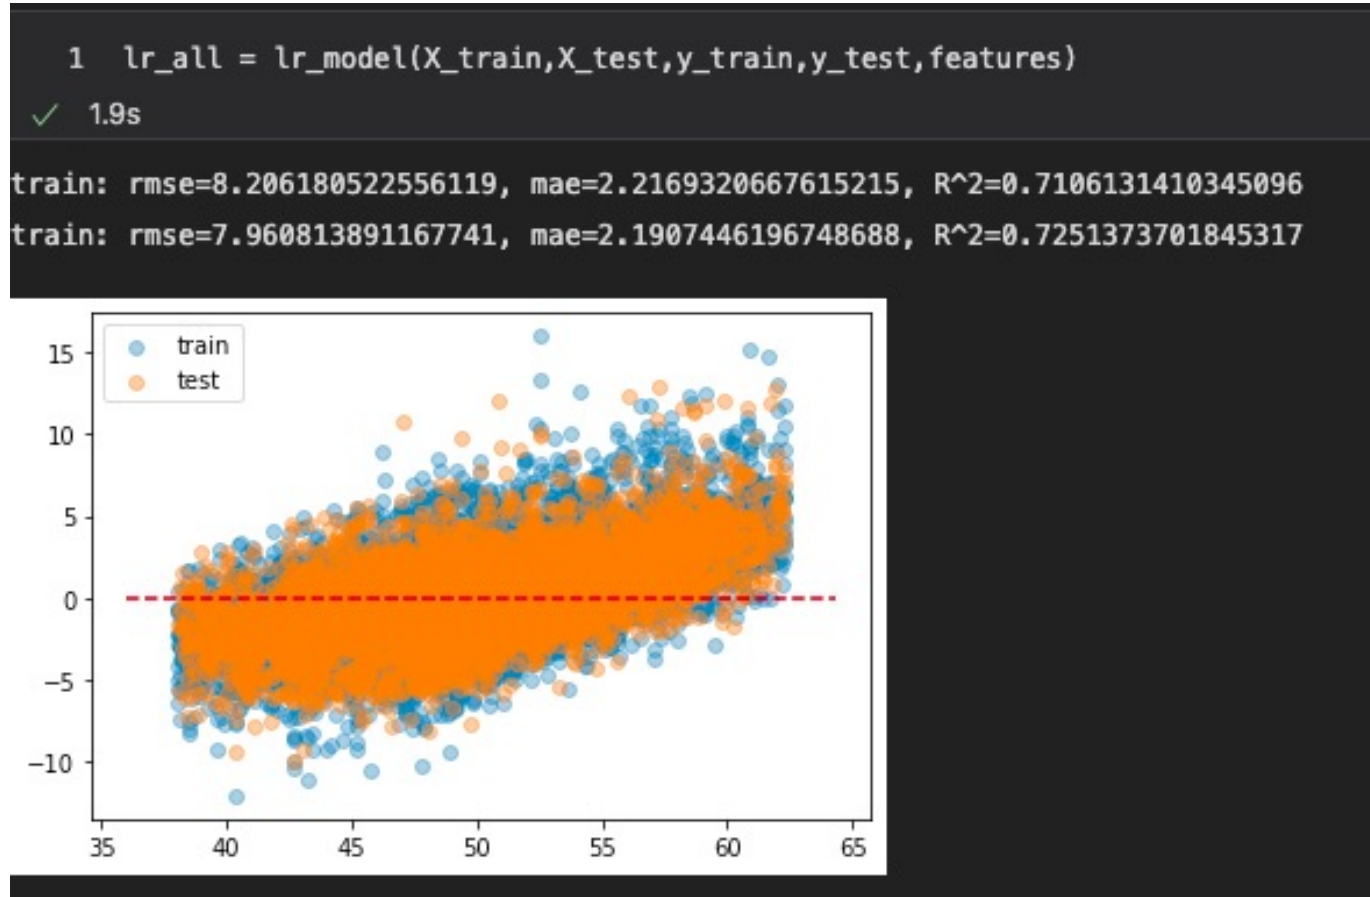
- Root Mean Squared Error (RMSE):

$$\sqrt{\frac{1}{N} \sum_i^N (t_i - y_i)^2} = \sqrt{\frac{RSS}{N}}$$

- Mean Squared Error (MSE):

$$\frac{1}{N} \sum_i^N |t_i - y_i|^2 = \frac{RSS}{N}$$

## 4.1.3. sklearn Linear Regression



	features	weight
9	total_rooms	-26.731592
13	population_per_household	-19.814847
3	INLAND	-5.188291
5	NEAR BAY	-1.619131
1	<1H OCEAN	-0.800810
6	NEAR OCEAN	-0.226635
12	housing_median_age	0.054975
4	total_bedrooms	0.411882
0	bedrooms_per_room	1.481176
2	population	6.031911
11	median_income	6.550643
7	ISLAND	7.834866
8	households	19.421038
10	rooms_per_household	19.867423


## 4.1.3. – sklearn Linear Regression

```
1 new_features = ['population_per_household', 'bedrooms_per_room', 'rooms_per_household', 'median_income', 'households',  
2 '<1H OCEAN', 'population', 'INLAND', 'NEAR BAY', 'NEAR OCEAN', 'ISLAND']  
3 lr_new = lr_model(X_train, X_test, y_train, y_test, new_features)
```

✓ 1.2s

train: rmse=8.547629583852357, mae=2.2610559882036503,  $R^2=0.6985721103657796$

train: rmse=8.324502107296473, mae=2.2414503816061435,  $R^2=0.71258032502751$





## 4.1. Linear Basis Function Models

4.1.1. Maximum Likelihood and Least Squares

4.1.2. Sequential Learning

4.1.3. sklearn Linear Regression Example

## 4.2. Regularization

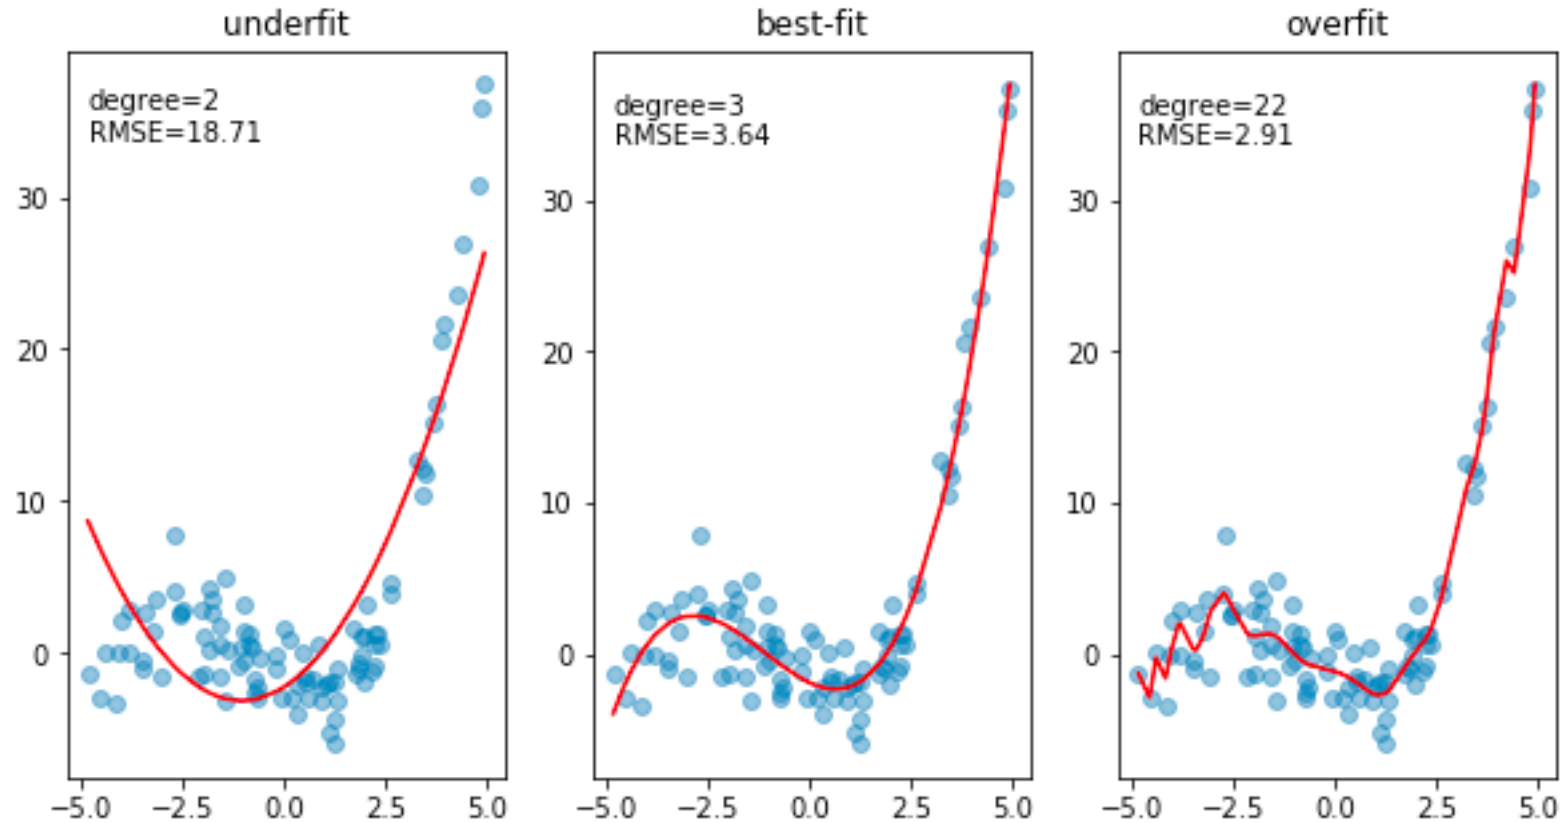
4.2.1. Overfit vs. Underfit

4.2.2. Regularized Least Squares

## 4.3. The Bias-Variance Decomposition

## 4.4. Conclusion

## 4.2.1. Overfit vs. Underfit





## 4.2.2. Regularized Least Squares

To control over-fitting, we can construct the total error function to be minimized in the form by adding the regularization term  $E_w(\mathbf{w})$

$$E_D(\mathbf{w}) + \lambda E_w(\mathbf{w})$$

where  $\lambda$  is the regularization coefficient that controls the relative importance of the data-dependent error  $E_D(\mathbf{w})$  and the regularization term  $E_w(\mathbf{w})$ .



## 4.2.2. Regularized Least Squares

Likelihood function  
See, slide 9.

Approaching from Bayesian theorem,

$$p(\mathbf{w}|D) = \frac{p(D|\mathbf{w})p(\mathbf{w})}{p(D)},$$

Prior probability of  $\mathbf{w}$

$p(D)$  is a normalizer so it can be ignored.

$$p(\mathbf{w}|D) = \prod_{n=1}^N p(D|\mathbf{w})p(\mathbf{w}) \rightarrow \ln p(\mathbf{w}|D) = \sum_{n=1}^N \ln p(D|\mathbf{w}) + \sum_{n=1}^N \ln p(\mathbf{w})$$

Assume the prior distribution of the coefficients are i.i.d. Gaussian,  $w_k \sim \mathcal{N}(0, 1/\lambda)$ , then

$$\sum p(\mathbf{w}) = \frac{\lambda}{2} \mathbf{w}^T \mathbf{w}.$$

The total error function becomes

$$E = E_D + \lambda E_w = \frac{1}{2} \sum_{n=1}^N \{t_n - \mathbf{w}^T \boldsymbol{\phi}\}^2 + \frac{1}{2} \lambda \mathbf{w}^T \mathbf{w}.$$


Solving for  $\mathbf{w}$ ,

$$\mathbf{w} = (\lambda \mathbf{I} + \boldsymbol{\Phi}^T \boldsymbol{\Phi})^{-1} \boldsymbol{\Phi}^T \mathbf{t}. \quad (4-13)$$

## 4.2.2. Regularized Least Squares

A more general regularized error form is


$$E = \frac{1}{2} \sum_{n=1}^N \{t_n - \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x}_n)\}^2 + \frac{\lambda}{2} \sum_{j=1}^M |w_j|^q. \quad (4-14)$$



Usually,  $q = 1|2$ .



## 4.2.2. Regularized Least Squares



If  $q = 2$ , it is called the *ridge* regularization.

- The function is strictly convex and differentiable everywhere.
- The dense solutions are possible, and all features are used.

If  $q = 1$ , it is called the *lasso* regularization.

- If  $\lambda$  is sufficiently large, some of the coefficients  $w_j$  are driven to zero.
- It leads to a *sparse* solution



## 4.2.2. Regularized Least Squares

Often, we can include them both regularizations. This is called *elastic net* regularization in the form

$$E = \frac{1}{2} \sum_{n=1}^N \{t_n - w^T \phi(x_n)\}^2 + \frac{\lambda}{2} \sum_{j=1}^M |w_j|^2 + |w_j|.$$

- The function is not differentiable.
- We have a unique solution.

## 4.2.2. Regularized Least Squares

```
1 from sklearn.linear_model import Lasso, Ridge, ElasticNet
```

✓ 0.1s

```
1 np.random.seed(123)
2 X = 10*np.random.sample(100)-5
3 X1 = np.array(sorted(X))
4 X2 = X**2
5 X3 = X**3
6 X4 = X**3/2
7 X = np.column_stack((X1,X2,X3,X4))
```

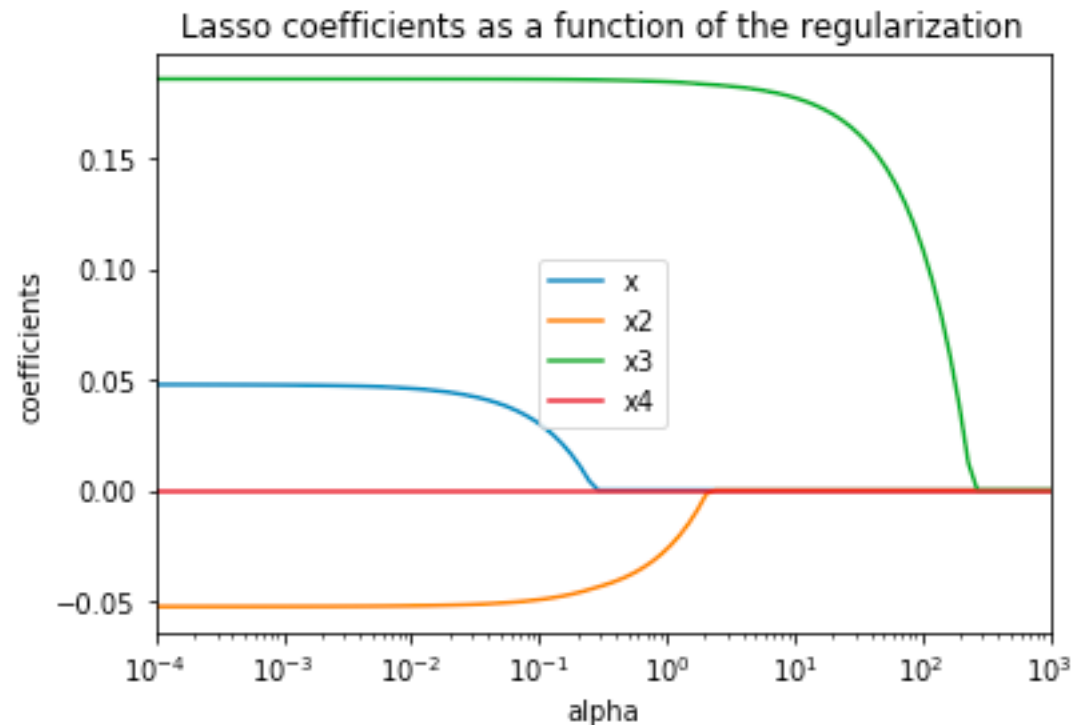
✓ 0.1s



## 4.2.2. Regularized Least Squares

```
1  alphas_lasso = np.logspace(-4, 3, 100)
2  coef_lasso = []
3  lasso = Lasso()
4  rmse, mae, rsq = [], [], []
5  for i in alphas_lasso:
6      lasso.set_params(alpha = i).fit(X,Y)
7      coef_lasso.append(lasso.coef_)
8      y_train_pred = lasso.predict(X)
9
10     rmse.append(mean_squared_error(Y,y_train_pred)) #Calculate root-mean-squared-error
11     mae.append(mean_absolute_error(Y,y_train_pred)) #Calculate mean-absolute-error
12     rsq.append(r2_score(Y,y_train_pred)) #Calculate R^2
13
14  features = ['x', 'x2', 'x3', 'x4']
15  df_coef = pd.DataFrame(coef_lasso, index=alphas_lasso, columns=features)
16  title = 'Lasso coefficients as a function of the regularization'
17  df_coef.plot(logx=True, title=title)
18  plt.xlabel('alpha')
19  plt.ylabel('coefficients')
20  plt.show()
✓ 2.2s
```


## 4.2.2. Regularized Least Squares



Lasso regularization therefore can be used for feature selection.

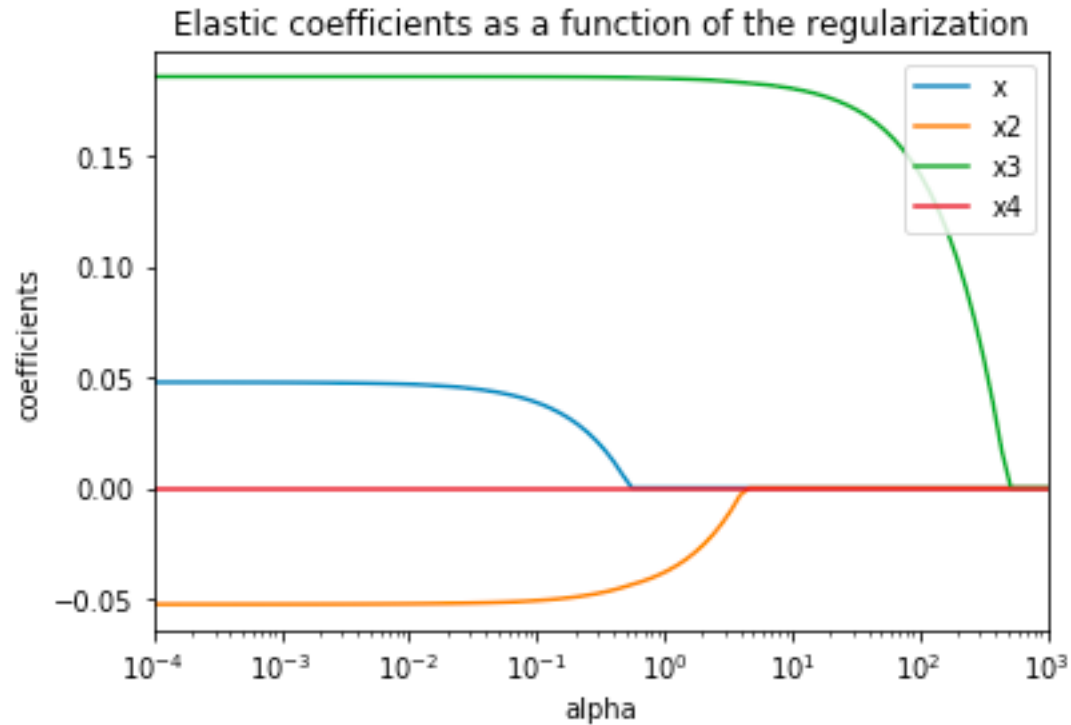
```
Lasso, alpha=0.01, intercept=0.24, coefficients=[ 0.04591006 -0.05221259  0.18560946  0.          ], rmse=13.006343902736498
Lasso, alpha=0.17783, intercept=0.21, coefficients=[ 0.01648733 -0.04703605  0.18533752  0.          ], rmse=13.012893764555317
Lasso, alpha=3.16228, intercept=-0.08, coefficients=[ 0.          -0.          0.18232391  0.          ], rmse=13.136403351426349
Lasso, alpha=56.23413, intercept=-0.05, coefficients=[ 0.          -0.          0.14271235  0.          ], rmse=15.48918780840939
Lasso, alpha=1000.0, intercept=0.04, coefficients=[-0.  0.  0.  0.], rmse=58.82740534381207
Linear Regression, intercept=0.24, coefficients=[ 0.04766321 -0.05252104  0.14850053  0.07425027], rmse=13.006323124645764
Linear Regression, intercept=0.24, coefficients=[ 0.04766321 -0.05252104  0.18562567], rmse=13.006323124645764
```

## 4.2.2. Regularized Least Squares



```
Ridge, alpha=0.01, intercept=0.17, coefficients=[ 0.0164893 -0.04107874  0.14741995  0.07370997], rmse=13.018155082313179
Ridge, alpha=0.17783, intercept=0.17, coefficients=[ 0.0164893 -0.04107874  0.14741995  0.07370997], rmse=13.018155082313179
Ridge, alpha=3.16228, intercept=0.17, coefficients=[ 0.0164893 -0.04107874  0.14741995  0.07370997], rmse=13.018155082313179
Ridge, alpha=56.23413, intercept=0.17, coefficients=[ 0.0164893 -0.04107874  0.14741995  0.07370997], rmse=13.018155082313179
Ridge, alpha=1000.0, intercept=0.17, coefficients=[ 0.0164893 -0.04107874  0.14741995  0.07370997], rmse=13.018155082313179
Linear Regression, intercept=0.24, coefficients=[ 0.04766321 -0.05252104  0.14850053  0.07425027], rmse=13.006323124645764
```

## 4.2.2. Regularized Least Squares



```
Elastic, alpha=0.01, intercept=0.24, coefficients=[ 0.04674513 -0.05235867  0.18561665  0.          ], rmse=13.006328837484837
Elastic, alpha=0.17783, intercept=0.22, coefficients=[ 0.03156567 -0.04964837  0.1854661  0.          ], rmse=13.008084572066696
Elastic, alpha=3.16228, intercept=-0.0, coefficients=[ 0.          -0.01232621  0.18348409  0.          ], rmse=13.084889385363315
Elastic, alpha=56.23413, intercept=-0.06, coefficients=[ 0.          -0.          0.16033351  0.          ], rmse=13.923383716994287
Elastic, alpha=1000.0, intercept=0.04, coefficients=[-0.  0.  0.  0.], rmse=58.82740534381207
Linear Regression, intercept=0.24, coefficients=[ 0.04766321 -0.05252104  0.14850053  0.07425027], rmse=13.006323124645764
```



## 4.1. Linear Basis Function Models

4.1.1. Maximum Likelihood and Least Squares

4.1.2. Sequential Learning

4.1.3. sklearn Linear Regression Example

## 4.2. Regularization

4.2.1. Overfit vs. Underfit

4.2.2. Regularized Least Squares

## 4.3. The Bias-Variance Decomposition

## 4.4. Conclusion



## 4.3. The Bias-Variance Decomposition


The average of expected loss also can be expressed as

$$\mathbb{E}(L) = \int \int (y - t)^2 p(x, t) dx dt.$$

(4-14)

If  $\mathbb{E}(L)$  is minimized, then  $\frac{\partial \mathbb{E}(L)}{\partial y} = 0$  solves  $y = \mathbb{E}_t[t|x]$ .

$\mathbb{E}_t(t|\mathbf{x})$  is the regression function that is the average over the ensemble of data sets and is the optimal prediction that is given by the conditional expectation.





## 4.3. The Bias-Variance Decomposition

We will use  $h = \mathbb{E}_t[t|x]$  for the simplicity.

When the error function  $(y - t)^2$  in Eq. 4-14 is expanded as

$$\begin{aligned}(y - t)^2 &= (y - h + h - t)^2 \\ &= (y - h)^2 + 2(y - h)(h - t) + (h - t)^2\end{aligned}\tag{4-15}$$

When Eq. (4-15) is substitute into Eq. (4-14), the second term vanishes and it becomes


$$\mathbb{E}(L) = \int \int (y - t)^2 p(x, t) dx dt = \int (y - h)^2 p(x) dx + \int \int (h - t)^2 p(x) dx dt\tag{4-16}$$

The first integration is the expectation of the data and it is the combination of  $(\text{bias})^2$  and variance.

$$\text{expected loss} = (\text{bias})^2 + \text{variance} + \text{noise}$$



# The Bias-Variance Decomposition



# The Bias-Variance Decomposition





Figure: Graphical illustration of bias vs. variance



## 4.1. Linear Basis Function Models

4.1.1. Maximum Likelihood and Least Squares

4.1.2. Sequential Learning

4.1.3. sklearn Linear Regression Example

## 4.2. Regularization

4.2.1. Overfit vs. Underfit

4.2.2. Regularized Least Squares

## 4.3. The Bias-Variance Decomposition

## 4.4. Conclusion



## 4.4. Conclusion

### 4.1. Linear Basis Function Models

- Assumptions:
  - Linearity:  $X$  and  $Y$  are in linear relationship.
  - Homoscedasticity: The variance of residual is the same for any features.
  - Independence: Examples are independent of each other.
  - Normality: All features are in Gaussian distribution.
- Optimization
  - GD, SGD, Mini-Batch GD

### 4.2. & 4.3. Regularization & Bias-Variance Decomposition

- The total errors are composite of Bias, Variance, and Noise.
  - Even though we want to optimize the model by reducing the errors,
    - Need to consider the trade-off between the bias and variance
    - Need to consider the trade-off between the overfitting and underfitting
    - Regularizations help to avoid the overfitting easily.
      - Lasso regularization also may be used for the feature selection or feature extraction.



## 4.4. Conclusion

Overall, The linear regression model implementation and interpretation are easy and straightforward. But,

- Simplicity: The model is too simple to capture the complexity of real data.
  - It is an easy start of modeling but hardly will be a final model.
  - It is essential to have different model.
- Assumptions are not realistic.
  - Need to do feature engineering, transformation, and more.
- Sensitivity: The model is sensitive to outliers as they are accounted in the estimation causing high variance and low bias.
  - Remove outliers.