



CS 559 – Machine Learning: Fundamental and Application

Lecture 6 – Neural Network

Outline



- 6.1. Lecture 5 Review
- 6.2. Introduction to Neural Network
- 6.3. Feedforward Functions
- 6.4. NN – Regression
- 6.5. NN – Classification
- 6.6. Optimization
- 6.7. Backpropagation
- 6.8. Numerical Example
- 6.9. Conclusion

6.1. Lecture 5 Review



In Lecture 5, we discussed linear classifiers.

1. The assumptions in linear regression will be carried over.
2. Discriminant Model – Directly assigns the target class label.
 1. LDA: Transforms into 1-D space and constructs the hyperplane model.
 2. Perceptron: Updates weights from the previous learn and constructs the hyperplane model.
3. Probabilistic Model – Assigns the target class label based on the Bayes learning
 1. Generative Model – MLE
 1. Finds the posterior probability of weights
 2. Discriminative Model – Logistic Regression
 1. Uses the logistic sigmoid function and stochastic descent optimization



6.1. Review - Optimization

Gradient Descent (GD)

$$\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla_{\mathbf{w}} L(\mathbf{w})$$

$\nabla_{\mathbf{w}} \text{TrainLoss}$ denotes the gradient of the (average) total training loss with respect to \mathbf{w} .

Stochastic Gradient Descent (SGD)

For each $(x, y) \in D_{\text{train}}$

$$\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla_{\mathbf{w}} L(x, y, \mathbf{w})$$

$\nabla_{\mathbf{w}} L$ denotes the gradient of one example loss with respect to \mathbf{w} .

6. Neural Networks



6.1. Lecture 5 Review

6.2. Introduction to Neural Network

6.3. Feedforward Functions

6.4. NN – Regression

6.5. NN – Classification

6.6. Optimization

6.7. Backpropagation

6.8. Numerical Example

6.9. Conclusion



6.2. Introduction to Neural Networks

- Often associated with biological devices (brains), electronic devices, or network diagrams
 - Proposed by Alexander Bain in 1873 and William James 1890
 - Interactions among neurons within the brain
- But the best conceptualization for this presentation is none of these: Think of neural network as a mathematical function.

6.2. Introduction to Neural Networks



Predicting Car Collision (Yes or No)

Suppose the 1-D coordinate of two vehicles (x_1 and x_2) are known. Predict if the collision will occur or not based on the simple classification model

$$y = \text{sign}(|x_1 - x_2| - 1).$$

We will say no if $y = +1$ and yes if $y = -1$.

Input: position of two oncoming cars $\mathbf{x} = [x_1, x_2]^T$

Output:

- $\mathbf{x} = [1, 3]^T, y = +1$; No
- $\mathbf{x} = [3, 1]^T, y = +1$; No
- $\mathbf{x} = [1, 0.5]^T, y = -1$; Yes

x_1	x_2	y
1	3	1
3	1	1
1	0.5	-1

6.2. Introduction to Neural Networks



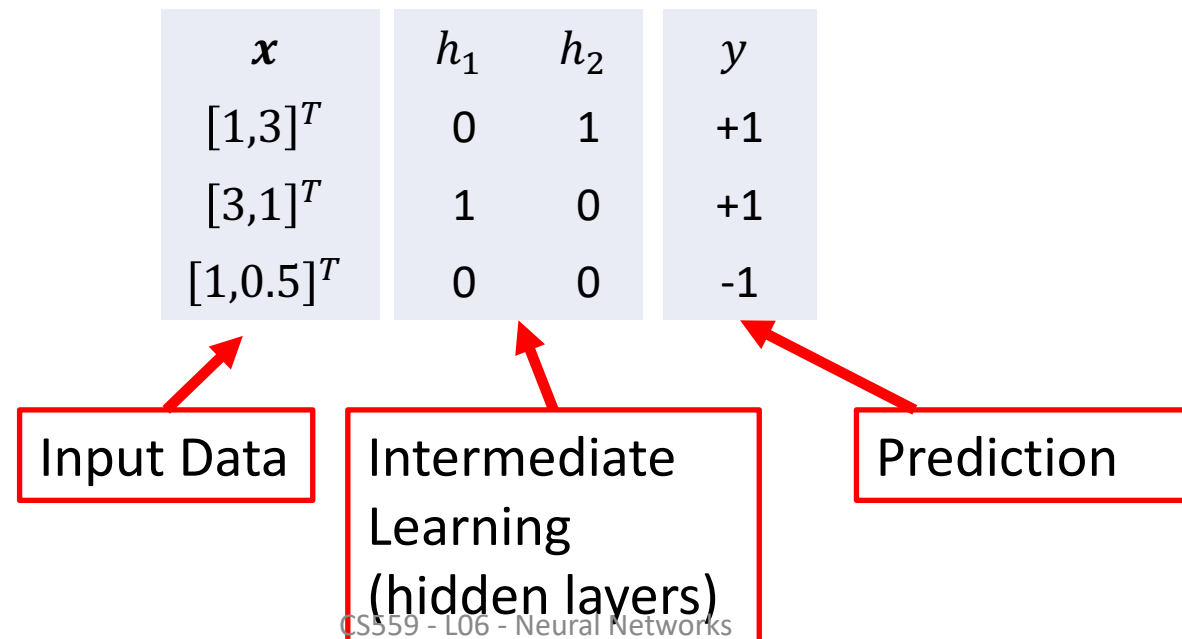
Predicting Car Collision (Yes or No)

Suppose we know when y will be +1 based on the position of one vehicle with respect to another vehicle. Let the new learning model be

$$y = \text{sign}(h_1 + h_2)$$

where h_i determines if the vehicle i is far right of j vehicle $h_i = (x_i - x_j \geq 1: 0)$.

The collision will not occur if either h is true.



6. Neural Networks



6.1. Lecture 5 Review

6.2. Introduction to Neural Network

6.3. Feedforward Network Functions

6.4. NN – Regression

6.5. NN – Classification

6.6. Optimization

6.7. Backpropagation

6.8. Numerical Example

6.9. Conclusion

6.3. Feedforward Network Functions





□ Consider a linear model based on linear combinations or fixed nonlinear basis function $\phi_j(\mathbf{x})$:

$$y(\mathbf{x}, \mathbf{w}) = f\left(\sum_{j=1}^M w_j \phi_j(\mathbf{x})\right) \quad (6-1)$$

where $f(\cdot)$ is a nonlinear activation function for classification and identity for regression.

- Recall that perceptron uses the activation function.

6.3. Feedforward Network Functions



6.3. Feedforward Network Functions



□ Neural networks (NN) use basis functions that follow the same form as Eq (1), so that each basis function is itself a nonlinear function of a linear combination of the inputs, where the coefficients in the linear combination are adaptive parameters.

□ The basic NN model can be described as a series of functional transformation:
1. Construct M linear combinations of input variables x_1, \dots, x_D in the form of

$$a_j = \sum_{i=1}^D w_{ji}^{(1)} x_i + w_{j0}^{(1)} \quad (6-2)$$

where $j = 1, \dots, M$, and the superscript (1) indicates that the corresponding parameters are in the first 'layer' of the network.

$w_{ji}^{(1)}$: weights
 $w_{j0}^{(1)}$: biases



6.3. Feedforward Network Functions

2. Transform each a_j using a differentiable, nonlinear activation function $h(\cdot)$,

$$z_j = h(a_j): \tag{6-3}$$

- These are called hidden units in NN that corresponds to the output $\phi_j(\mathbf{x})$.
- The nonlinear $h(\cdot)$ are either the logistic sigmoid or *tanh* function.

3. Linearly recombine to give output unit activations:

$$a_k = \sum_{j=1}^M w_{kj}^{(2)} z_j + w_{k0}^{(2)} \tag{6-4}$$

where $k = 1, \dots, K$ and K is the total number of outputs.

4. Transform the output unit activations (6-4) to predict y_k :

- For standard linear regression: $y_k = a_k$
- For binary classification: $y_k = \sigma(a_k)$
- For multiclassification: softmax



6.3. Feedforward Network Functions

- ❑ The NN **forward propagation** can be summarized as

$$y_k(x, w) = \sigma \left(\sum_{j=1}^M w_{kj}^{(2)} h \left(\sum_{i=1}^D w_{ji}^{(1)} x_i + w_{j0}^{(1)} \right) + w_{k0}^{(2)} \right) \quad (6-5)$$

- ❑ The NN model is simply a nonlinear function from a set of input variables $\{x_i\}$ to a set of output variables $\{y_k\}$ controlled by a vector w of adjustable parameters.
- ❑ We can absorb the bias into the set of weights as $a_j = \sum_{i=0}^D w_{ji}^{(1)} x_i$ then Eq (6-5) becomes

$$y_k(x, w) = \sigma \left(\sum_{j=0}^M w_{kj}^{(2)} h \left(\sum_{i=0}^D w_{ji}^{(1)} x_i \right) \right). \quad (6-6)$$

- ❑ NN is also known as the *multilayer perceptron* using continuous sigmoidal nonlinearities.
- ❑ If the activation function of the hidden units is linear, we always find an equivalent network.

6. Neural Networks



6.1. Lecture 5 Review

6.2. Introduction to Neural Network

6.3. Feedforward Network Functions

6.4. NN – Regression

6.5. NN – Classification

6.6. Optimization

6.7. Backpropagation

6.8. Numerical Example

6.9. Conclusion



6.4 NN – Regression

- ❑ Consider a set of input vectors $\{\mathbf{x}_n\}$ and target vectors $\{\mathbf{t}_n\}$ where $n = 1, \dots, N$.
- ❑ For the simplicity, assume that t has a Gaussian distribution with an \mathbf{x} -dependent mean

$$p(t|\mathbf{x}, \mathbf{w}) = \mathcal{N}(t|y(\mathbf{x}, \mathbf{w}), \beta^{-1})$$

where β is an inverse variance of the Gaussian noise.

- To take the output unit activation function be the identity.

- ❑ The corresponding likelihood function is

$$p(\mathbf{t}|\mathbf{X}, \mathbf{w}, \beta) = \prod_{n=1}^N p(t_n|\mathbf{x}_n, \mathbf{w}, \beta) \quad (6-7)$$

- ❑ We can learn the parameters \mathbf{w} and β by taking the negative log:

$$-\log p(\mathbf{t}|\mathbf{X}, \mathbf{w}, \beta) = \frac{\beta}{2} \sum_{n=1}^N \{y(\mathbf{x}_n, \mathbf{w}) - t_n\}^2 - \frac{N}{2} \ln \beta + \frac{N}{2} \ln(2\pi). \quad (6-8)$$



6.4. NN – Regression

□ In NN, it is usual to consider the minimization of an error function

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(\mathbf{x}_n, \mathbf{w}) - t_n\}^2 . \quad (6-8)$$

□ The maximum likelihood (ML) solution \mathbf{w}_{ML} is found by minimizing $E(\mathbf{w})$.



6.4. NN – Regression

- Using \mathbf{w}_{ML} , the value of β_{ML} can be found by minimizing Equation (6-8):

$$\frac{1}{\beta_{ML}} = \frac{1}{N} \sum_{n=1}^N \{y(\mathbf{x}_n, \mathbf{w}_{ML}) - t_n\}^2 \quad (6-9)$$

- If we have multiple target variables,

$$\frac{1}{\beta_{ML}} = \frac{1}{NK} \sum_{n=1}^N |y(\mathbf{x}_n, \mathbf{w}_{ML}) - \mathbf{t}_n|^2 \quad (6-10)$$

Where K is the number of target variables.

- An output activation is the identity, so $y_k = a_k$ and the corresponding error function has the property

$$\frac{\partial E}{\partial a_k} = y_k - t_k \quad (6-11)$$

6. Neural Networks



- 6.1. Lecture 5 Review
- 6.2. Introduction to Neural Network
- 6.3. Feedforward Network Functions
- 6.4. NN – Regression
- 6.5. NN – Classification**
- 6.6. Optimization
- 6.7. Backpropagation
- 6.8. Numerical Example
- 6.9. Conclusion



6.5. NN – Classification

❑ Binary classification having a single target variable $t, t \in \{0,1\}$

❑ An activation function is a logistic sigmoid:

$$y = \sigma(a) = \frac{1}{1 + \exp(-a)}$$

so $0 \leq y(\mathbf{x}, \mathbf{w}) \leq 1$.

❑ We can interpret the conditional probability for each class

$$C_1: p(C_1|\mathbf{x}) = y(\mathbf{x}, \mathbf{w})$$

$$C_2: p(C_2|\mathbf{x}) = 1 - p(C_1|\mathbf{x})$$

And the conditional distribution $p(t|\mathbf{x}, \mathbf{w})$ is then a Bernoulli distribution:

$$p(t|\mathbf{x}, \mathbf{w}) = y(\mathbf{x}, \mathbf{w})^t \{1 - y(\mathbf{x}, \mathbf{w})\}^{1-t}$$

6.5. NN – Classification



- The cross-entropy error function after taking the negative log likelihood is then

$$E(w) = - \sum_{n=1}^N \{t_n \ln y_n + (1 - t_n) \ln(1 - y_n)\}. \quad (6-12)$$

Note $\beta = 0$ if the target values are assumed to be correctly labelled.

- In NN, training is faster and easier to generalize using the cross-entropy function.



6.5. NN – Classification

- ❑ K classes and $t_k \in \{0,1\}$ have a 1-of- K .
- ❑ The activation function is the *softmax* function:

$$y_k(\mathbf{x}, \mathbf{w}) = \frac{\exp(a_k(\mathbf{x}, \mathbf{w}))}{\sum_j \exp(a_j(\mathbf{x}, \mathbf{w}))} \quad (6-13)$$

where $0 \leq y_k \leq 1$ and $\sum_k y_k = 1$.

- ❑ NN output $y_k(\mathbf{x}, \mathbf{w}) = p(t_k = 1|\mathbf{x})$ leads to the error function is the multiple cross-entropy function

$$E(\mathbf{w}) = - \sum_{n=1}^N \sum_{k=1}^K t_{kn} \ln y_k(\mathbf{x}_n, \mathbf{w}) \quad (6-14)$$

6. Neural Networks



- 6.1. Lecture 5 Review
- 6.2. Introduction to Neural Network
- 6.3. Feedforward Network Functions
- 6.4. NN – Regression
- 6.5. NN – Classification
- 6.6. Optimization**
- 6.7. Backpropagation
- 6.8. Numerical Example
- 6.9. Conclusion

6.6. Optimization


- ❑ Use Gradient descent optimization by updating weight

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} - \eta \nabla E(\mathbf{w}^{(\tau)})$$

where η is the learning rate being > 0 and τ is the labelled iteration number.

- ❑ To evaluate the gradient of error efficiently, we use **error backpropagation** through the network.

- ❑ In iteration $[t + 1]$, choose a weight update $\Delta w[t]$ and set $w[t + 1] = w[t] + \Delta w[t]$.




6.6. Optimization



- ❑ Example: GD minimizes the error/loss $L_n(w)$ by taking steps in the direction of the negative gradient:

$$\Delta w[t] = -\eta \frac{\partial L_n}{\partial w[t]} \quad (6-15)$$

- ❑ Problem: How to evaluate $\frac{\partial L_n}{\partial w[t]}$ in iteration $[t + 1]$?
 - “Error Backpropagation” allows to evaluate $\frac{\partial L_n}{\partial w[t]}$ in $\mathcal{O}(W)$!




6. Neural Networks



- 6.1. Lecture 5 Review
- 6.2. Introduction to Neural Network
- 6.3. Feedforward Network Functions
- 6.4. NN – Regression
- 6.5. NN – Classification
- 6.6. Optimization
- 6.7. Error Backpropagation**
- 6.8. Numerical Example
- 6.9. Conclusion

6.7. Error Backpropagation

- ❑ Backpropagation is used to describe the training using gradient descent applied to an SSE function.
- ❑ Two steps:
 - ❑ Evaluation of derivatives.
 - ❑ Adjustment of weights.



6.7. Error Backpropagation

□ Consider a *forward propagation*:

□ Let the error function be $E(\mathbf{w})$ and y_k be a linear combination of x_i :

$$y_k = \sum_i w_{ki} x_i$$

$$E_n = \frac{1}{2} \sum_k (y_{nk} - t_{nk})^2$$

□ The gradient of error r.t. w_{ji} - the "local" computation with product of an error associated with the output end of the link w_{ji} and x_{ni}

$$\frac{\partial E_n}{\partial w_{ji}} = (y_{nj} - t_{nj}) x_{ni} = \frac{\partial E_n}{\partial a_j} \left(\frac{\partial a_j}{\partial w_{ji}} \right) = \delta_j z_i \quad (6-16)$$

$$\delta_j = \frac{\partial E_n}{\partial a_j} = \sum_k \frac{\partial E_n}{\partial a_k} \left(\frac{\partial a_k}{\partial a_j} \right)$$

$$\frac{\partial a_j}{\partial w_{ji}} = z_i \text{ and } z_j = h(a_j)$$

$$a_j = \sum_i w_{ji} z_i$$


6.7. Error Backpropagation

□ We can obtain the backpropagation formula by

$$\delta_j = h'(a_j) \sum_k w_{kj} \delta_k \quad (6-17)$$

Derivation:

$$\delta_j = \sum_k \frac{\partial E_n}{\partial a_k} \frac{\partial a_k}{\partial a_j} = \sum_k \delta_k \frac{\partial}{\partial a_j} \left(\sum_i w_{ki} z_i \right) = \sum_k \delta_k \frac{\partial}{\partial a_j} \left(\sum_i w_{ki} h(a_i) \right)$$




$$\begin{aligned} \delta_k &= \frac{\partial E_n}{\partial a_k} \\ a_k &= \sum_i w_{ki} z_i \\ z_i &= h(a_i) \end{aligned}$$

All partial derivatives are 0 except when $j = i$

6.7. Error Backpropagation

□ Procedure Summary:

1. Apply an input x_n to the network and forward propagate through network using definitions in Eq. (6-16) to find the activations of all the hidden and output units.
2. Evaluate δ_k for all the output units : $\delta_k = y_k - t_k$.
3. Backpropagate and obtain δ_j for each hidden unit in the network using Equation (6-17).
4. Use Equation (6-16) to evaluate the required derivatives.



6.7. Error Backpropagation



□ Example: two-layer network regression ($y_k = a_k$) with hidden units $h(a) = \tanh(a) = \frac{e^a - e^{-a}}{e^a + e^{-a}}$.

□ $E_n = \frac{1}{2} \sum_{k=1}^K (y_k - t_k)^2$

□ $h'(a) = \frac{dh}{da} = 1 - h(a)^2$

□ A forward propagation using $a_j = \sum_{i=0}^D w_{ji}^{(1)} x_i$, $z_j = \tanh(a_j)$, $y_k = \sum_{j=0}^M w_{kj}^{(2)} z_j$

□ Compute δ 's for each output unit using $\delta_k = y_k - t_k$

□ Then backpropagate to obtain δ s for the hidden units using

$$\delta_j = (1 - z_j^2) \sum_{k=1}^K w_{kj} \delta_k$$

□ Derivatives w.r.t. the first-layer and second-layer weights:

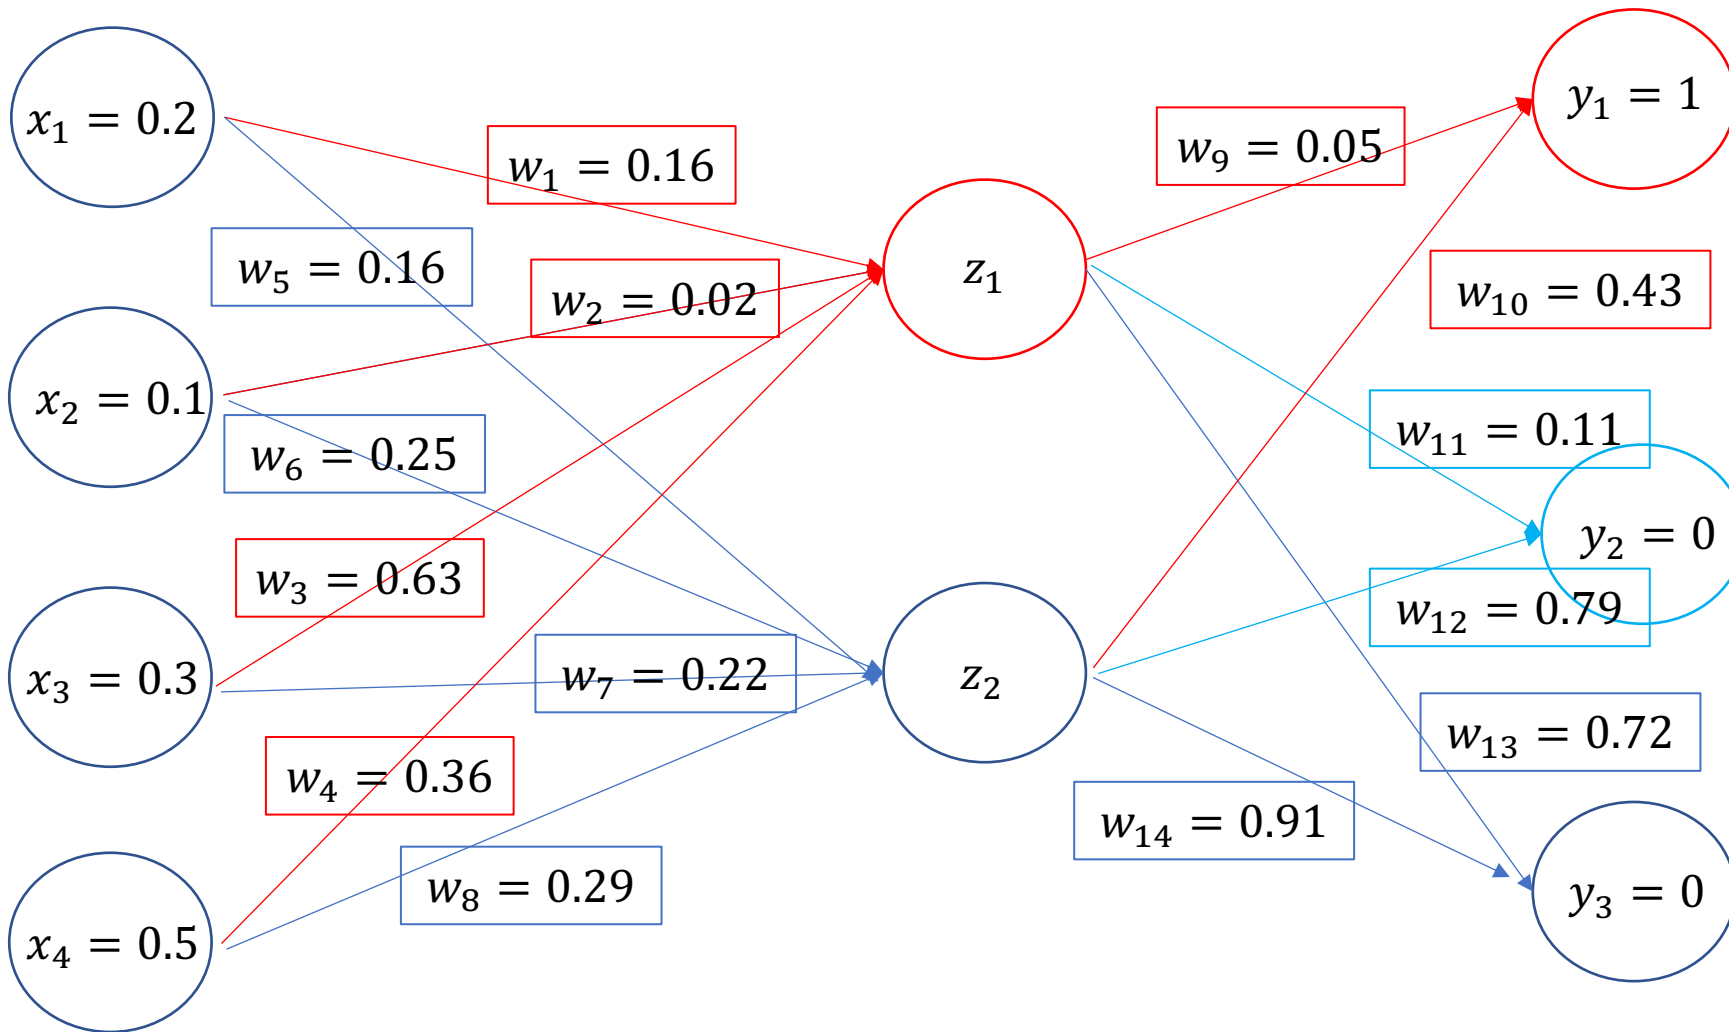
$$\frac{\partial E_n}{\partial w_{ji}^{(1)}} = \delta_j x_i, \frac{\partial E_n}{\partial w_{kj}^{(2)}} = \delta_k z_j$$

6. Neural Networks



- 6.1. Lecture 5 Review
- 6.2. Introduction to Neural Network
- 6.3. Feedforward Network Functions
- 6.4. NN – Regression
- 6.5. NN – Classification
- 6.6. Optimization
- 6.7. Error Backpropagation
- 6.8. Numerical Example**
- 6.9. Conclusion

6.8. Numerical Example



6.8. Numerical Example



Forward Propagation using Eq. 6-3:

$$\mathbf{z} = \sigma(\mathbf{w}_1^T \mathbf{x}) = \frac{1}{1 + \exp(-\mathbf{w}_1^T \mathbf{x})} = \frac{1}{1 + \exp\left(-\left([0.2 \ 0.1 \ 0.5] \begin{bmatrix} 0.16 & 0.16 \\ 0.02 & 0.25 \\ 0.63 & 0.22 \\ 0.36 & 0.29 \end{bmatrix}\right)\right)} = [0.5994 \ 0.5666]$$

Prediction using Eq. 6-5: $\hat{\mathbf{y}} = \sigma(\mathbf{w}_2^T \mathbf{z})$

$$\hat{\mathbf{y}} = \frac{1}{1 + \exp(-\mathbf{w}_2^T \mathbf{z})} = \frac{1}{1 + \exp\left(-\left([0.5994 \ 0.5666] \begin{bmatrix} 0.05 & 0.33 & 0.72 \\ 0.43 & 0.79 & 0.91 \end{bmatrix}\right)\right)} = [0.5680 \ 0.6560 \ 0.7205]$$

Error Function using Eq. 6-6:
$$L = \sum_{i=1} \frac{1}{2} (\hat{\mathbf{y}} - \mathbf{y})^2 = 0.5681$$

6.8. Numerical Example



Back Propagation from output to last hidden layer: $\mathbf{G}_p = \mathbf{E}'\hat{\mathbf{y}}'$

Derivative of error function (Eq. 6-11):

$$\mathbf{E}' = \hat{\mathbf{y}} - \mathbf{y} = [0.5680 \ 0.6560 \ 0.7205] - [1 \ 0 \ 0] = [-0.4320 \ 0.6560 \ 0.7205]$$

$$\text{Derivative of } \hat{\mathbf{y}}: \hat{\mathbf{y}}' = \hat{\mathbf{y}}(1 - \hat{\mathbf{y}}) = [0.2454 \ 0.2257 \ 0.2014]$$

$$\text{The gradient of layer: } \mathbf{G}_p = \mathbf{E}'\hat{\mathbf{y}}' = [-0.1060 \ 0.1480 \ 0.1451]$$

$$\delta \mathbf{W}_2 = (\mathbf{G}_p^T \mathbf{z})^T = \left(\begin{bmatrix} -0.1060 \\ 0.1480 \\ 0.1451 \end{bmatrix} [0.5994 \ 0.5666] \right)^T = \begin{bmatrix} -0.0635 & 0.0887 & 0.0870 \\ -0.0601 & 0.0839 & 0.0822 \end{bmatrix}$$

$$\widehat{\mathbf{W}}_2 = \mathbf{w}_2 - \delta \mathbf{W}_2 = \begin{bmatrix} 0.05 & 0.33 & 0.72 \\ 0.43 & 0.79 & 0.91 \end{bmatrix} - \begin{bmatrix} -0.0635 & 0.0887 & 0.0870 \\ -0.0601 & 0.0839 & 0.0822 \end{bmatrix} = \begin{bmatrix} 0.1135 & 0.2413 & 0.6330 \\ 0.4901 & 0.7061 & 0.8278 \end{bmatrix}$$

6.8. Numerical Example



Back Propagation from last hidden layer

Derivative of z : $\mathbf{z}' = \mathbf{z}(1 - \mathbf{z}) = [0.5994(1 - 0.5994) \ 0.5666(1 - 0.5666)] = [0.2401 \ 0.2456]$

The gradient of 1st layer: $\mathbf{G} = \mathbf{G}_p \mathbf{w}_2^T \mathbf{z}'$

$$\mathbf{G} = [-0.1060 \ 0.1480 \ 0.1451] \begin{bmatrix} 0.05 & 0.43 \\ 0.33 & 0.79 \\ 0.72 & 0.91 \end{bmatrix} [0.2401 \ 0.2456] = [0.0355 \ 0.0499]$$

$$\delta \mathbf{W}_1 = \mathbf{x}^T \mathbf{G} = \begin{bmatrix} 0.2 \\ 0.1 \\ 0.3 \\ 0.5 \end{bmatrix} [0.0355 \ 0.0499] = \begin{bmatrix} 0.0071 & 0.0100 \\ 0.0036 & 0.0050 \\ 0.0107 & 0.0150 \\ 0.0178 & 0.0250 \end{bmatrix}$$

$$\widehat{\mathbf{W}}_1 = \mathbf{w}_1 - \delta \mathbf{W}_1 = \begin{bmatrix} 0.16 & 0.16 \\ 0.02 & 0.25 \\ 0.63 & 0.22 \\ 0.36 & 0.29 \end{bmatrix} - \begin{bmatrix} 0.0071 & 0.0100 \\ 0.0036 & 0.0050 \\ 0.0107 & 0.0150 \\ 0.0178 & 0.0250 \end{bmatrix} = \begin{bmatrix} 0.1529 & 0.1500 \\ 0.0164 & 0.2450 \\ 0.6193 & 0.2050 \\ 0.3422 & 0.2650 \end{bmatrix}$$

6.8. Numerical Example



Calculate new z :

$$z' = \frac{1}{1 + \exp(-\hat{W}_1^T x)} = [0.5961 \ 0.5618]$$

Predict y :

$$\hat{y} = \frac{1}{1 + \exp(-\hat{W}_2^T z)} = [0.5849 \ 0.6319 \ 0.6990]$$

Estimate SSE:

$$E = \frac{1}{2} \sum (\hat{y} - y)^2 = 0.5301$$

6. Neural Networks



- 6.1. Lecture 5 Review
- 6.2. Introduction to Neural Network
- 6.3. Feedforward Network Functions
- 6.4. NN – Regression
- 6.5. NN – Classification
- 6.6. Optimization
- 6.7. Backpropagation
- 6.8. Numerical Example
- 6.9. Conclusion**

6.9. Conclusion



Pros:

- Successfully used on a **variety of domains**: computer vision, speech recognition, gaming, etc.
- Can provide solutions to very **complex and nonlinear** problems.
- If provided with **sufficient amount of data**, can solve classification and forecasting problems accurately and easily.
- Once trained, **prediction is fast**.

Cons:

- The outcome contains some uncertainty that is not always desirable.
- It has to undergo a “learning” phase.
- The quality of the outcome depends on the quality of the data used in the learning phase.